

MediAssist AI

Healthcare Workflow Support System

AI for Bharat Hackathon 2026 | Healthcare & Life Sciences Track



Team Cyber Scorpion

Team Lead: Bhavik Patel

Team Member: Hem Bhavsar

March 2026

1. Executive Summary

MediAssist AI is an AI-powered clinical document processing prototype that addresses a critical bottleneck in healthcare — the time clinicians spend translating complex medical records into patient-friendly language. Built on AWS serverless architecture with a local LLM, MediAssist AI provides:

- Automated professional clinical summaries from uploaded documents
- Plain-language patient explanations generated by Ollama LLM
- Human-in-the-loop approval workflow ensuring doctor oversight
- Full audit trail for regulatory compliance
- Support for PDF, DOCX, images (JPG/PNG), and scanned documents via OCR

 **Important:** This system processes synthetic data only and is a demonstration prototype. It does not provide medical diagnoses or treatment recommendations.

2. Problem Statement

The Challenge

Healthcare professionals in India and globally face a persistent problem: complex clinical documentation is difficult for patients to understand, leading to poor health literacy, medication non-compliance, and worse outcomes.

Challenge	Impact
Clinical language is highly technical	Patients cannot understand their own health records
Doctors lack time for detailed explanations	Average consultation: 2–5 minutes
Manual summarization is error-prone	Risk of miscommunication and harm
No audit trail for AI-generated content	Compliance and liability concerns
Scanned/image documents are inaccessible	Large volume of legacy paper records in India

The Opportunity

India has 1 doctor per 834 patients (WHO standard: 1 per 1000). AI-assisted documentation can help stretch clinical capacity while keeping humans firmly in control of medical decisions.

3. Our Solution

MediAssist AI provides a 4-step workflow for clinical document processing:


Step	Description
① Upload	Clinician uploads PDF, DOCX, image or text document
② Process	Document sent to AWS Lambda, text extracted, AI triggered async
③ Review	Doctor reviews AI summary, edits if needed, approves release
④ Export	Approved explanation exported with full audit trail

Key Differentiators

- No API timeout issues — async worker pattern handles Ollama's 60-90s processing time
 - Privacy-first — LLM runs on private EC2, no data sent to third-party APIs
 - OCR support — Amazon Textract handles scanned PDFs and photos of documents
 - Human oversight — AI cannot release content without doctor approval
 - Compliance-ready — every action logged with timestamp and reviewer ID
-

4. System Architecture

Architecture Diagram



Why Async Pattern?

Problem: AWS API Gateway has a hard 29-second timeout. Ollama LLM takes 60-90 seconds to process clinical documents.

Solution: Process Lambda returns doc_id instantly (<1s). Worker Lambda runs Ollama in background (600s timeout). React frontend polls /status every 5 seconds.

AWS Services Used

Service	Purpose
AWS Lambda (×5)	Serverless compute — Process, Worker, Status, Approve, Audit
Amazon API Gateway	HTTP API with CORS — 4 routes
Amazon DynamoDB	Results, Users, AuditLog tables (PAY_PER_REQUEST)
Amazon S3	Document storage + React frontend hosting
Amazon Textract	OCR for scanned PDFs and image uploads
Amazon CloudWatch	Live monitoring dashboard for all Lambda functions
AWS IAM	Least-privilege roles for all services

Lambda Functions

Function	Timeout / Memory / Role
MediAssist-Process	30s / 512MB — Receives upload, extracts text, triggers Worker async
MediAssist-Worker	600s / 512MB — Calls Ollama, updates DynamoDB with AI result
MediAssist-Status	30s / 256MB — Polled by React every 5s to check completion
MediAssist-Approve	30s / 256MB — Doctor approves/edits summary, updates status
MediAssist-Audit	30s / 256MB — Returns full compliance audit log

5. Technology Stack

Frontend

Technology	Usage
React 19.2 + Vite	Single-page application with fast HMR development
TailwindCSS	Responsive UI styling
React Context API	Global state management across 4 workflow panels
Fetch API	REST calls + polling pattern for async status

Backend

Technology	Usage
Python 3.11	All Lambda function runtime
PyPDF2	Text extraction from PDF files
python-docx	Text extraction from DOCX files
boto3	AWS SDK — DynamoDB, S3, Lambda, Textract
Ollama + llama3.2:3b	Local LLM on EC2 t3.large (8GB RAM)

Document Processing Support

File Type	Extraction Method
.txt	Direct UTF-8 decode
.pdf (text-based)	PyPDF2 page extraction
.pdf (scanned)	Auto-fallback to Amazon Textract OCR
.docx	python-docx paragraph extraction
.jpg / .png / .tiff	Amazon Textract OCR
.bmp / .webp	Amazon Textract OCR

6. Responsible AI Approach

Principle	How We Implement It
Human Oversight	Every AI output requires explicit doctor approval before patient sees it
Transparency	Confidence scores (High/Medium/Low) shown for every summary
Safety Constraints	LLM prompt explicitly prevents diagnostic/treatment recommendations
Audit Trail	Every action (upload, generate, edit, approve) logged with timestamp
Data Privacy	LLM runs on private EC2 — no clinical data sent to third-party APIs
Bias Awareness	Synthetic data only — no real patient data processed in prototype

Core Principle: The AI assists, the doctor decides. MediAssist AI cannot release any content to patients without explicit human approval.

7. API Reference

Endpoint	Description
POST /process	Upload document → returns {doc_id, status: PROCESSING} instantly
GET /status?doc_id=xxx	Poll for AI result — returns PROCESSING / PENDING / FAILED
POST /approve	Doctor approves/edits summary → status becomes APPROVED
GET /audit?doc_id=xxx	Returns full chronological audit log for compliance

API Gateway

Property	Value
Type	HTTP API v2 (not REST API)
CORS	AllowOrigins: *, Methods: GET, POST, OPTIONS
Region	ap-south-1 (Mumbai)
Base URL	https://68nnaciiv6.execute-api.ap-south-1.amazonaws.com/prod
Auth	None (demo) — add Cognito for production

8. Deployment

Infrastructure (AWS — ap-south-1 Mumbai)

Resource	Details
EC2 Instance	t3.large, 8GB RAM — running Ollama + llama3.2:3b
S3 Buckets	mediassistai-documents (storage), mediassistai-frontend (hosting)
DynamoDB Tables	MediAssist-Results, MediAssist-Users, MediAssist-AuditLog
IAM Role	MediAssist-Lambda-Role with S3, DynamoDB, Textract, Lambda invoke

Deployment Commands

```
# Full deployment
cd PrototypeDemo/infra && bash deploy.sh
# CloudWatch dashboard
bash setup_dashboard.sh
# Pre-warm Ollama before demo
curl http://3.109.55.138:11434/api/generate -d
'{"model":"llama3.2:3b","prompt":"hello","stream":false}'
```

9. Limitations & Future Work

Current Limitations (Prototype)

- No user authentication — add AWS Cognito for production
- English language only — multilingual support needed for AI for Bharat
- Processing time: 60-90 seconds — upgrade EC2 to GPU instance for <10s
- Scanned image quality affects OCR accuracy
- Not HIPAA certified — prototype uses synthetic data only

Roadmap for Production

Feature	Priority
Multilingual support (Hindi, Tamil, Gujarati)	High
GPU EC2 for faster inference	High
AWS Cognito authentication	High
HL7 FHIR integration	Medium
Mobile app (React Native)	Medium
Real-time collaboration for doctors	Low

10. Team

Detail	Information
Team Name	Cyber Scorpion 🦂
Team Lead	Bhavik Patel
Team Member	Hem Bhavsar
Hackathon	AI for Bharat 2026 (Powered by AWS)
Track	Healthcare & Life Sciences
Submission	March 2026
GitHub	github.com/BhavikPatel/AI-for-Healthcare-Life-Sciences

⚠️ This system uses synthetic data only. Not for clinical use.